


Training regime decides the width gap (121M, 3 seeds)



One checkpoint, three inference budgets

